Notes on the Julia Programming Language

Lutz Hendricks

UNC Chapel Hill

August 25, 2019

Abstract

This document summarizes my experience with the Julia language. Its main purpose is to document tips and tricks that are not covered in the official documentation.

1 My Setup (1.1)[1]

Since things are in flux, I find it useful to use the offical `JuliaPro` installation.

My startup file loads the packages `OhMyREPL` and `Revise`. `Revise` comes after packages from the standard libraries, so it does not track changes to those.

1.1 JuliaPro (1.1)

Sometimes JuliaPro gets slow and has trouble updating the REPL screen. Then restarting the computer is the only solution.

1.2 Julia + Editor (1.1)

It appears that the default editor is determined by the system wide file association. No need to set the `JULIA_EDITOR` environment variable.

---

[1] Each section is labeled with the Julia version for which it was last updated.

One drawback: Links in the terminal REPL are not clickable. A substantial drawback during debugging. So I end up using BBEdit as my main editor, but do some debugging in `Juno`. Not ideal.

2 Arrays (1.1)

2.1 Indexing

Extracting specific elements with indices given by vectors:

```
A = rand(4,3,5);
A[CartesianIndex.([1,2], [2,2]), 1] -> A[1,2,1] and A[2,2,1]
```

Similar to using sub2ind:

```
idxV = sub2ind(size(A), [1,2],[2,2],[1,1])
A[idxV]
```

To extract a "row" of a multidimensional matrix without hard-coding the dimensions, generate a `view` using `selectdim`.

3 Data Handling (1.1)

`DataFrames` has the rough equivalent of Matlab `table`s.

`JuliaDB` resembles the data handling capabilities of traditional stats packages (like STATA) where processing happens outside of RAM.

3.1 Dataframes (1.1)

Stats packages often provide data in "long" format (i.e., each id occupies many rows to cover many variables). unstack and `stack` can go back and forth between "long" and "short" formats.

4 Debugging (1.1)

The `Juno` debugger stopped working in V.1.1 (invoking it hangs `Julia`). But the command line debugger may well be the better option. After `using Debugger` invoke `@enter foo(x)` to start a debugging session.

Particularly useful:

- `bp on error`

- `bp add func:line` with possible restrictions on particular argument types.

5 External Programs (1.1)

One can execute `bash` commands with `run`.

Question: Trying to run a bash script using `run(`. myscript.sh`)` produces a permission denied error (even though permissions are set so that others can execute. Why?

6 Formatting (1.1)

The `Formatting` package seems to be the best bet. It uses `Python` like syntax and can format multiple arguments simultaneously (not well documented). Example:

```
fs = FormatExpr("{1:.2f} and {2:.3f}")
format(fs, 1.123, 2)
```

yields `"1.12 and 2.000"`.

7 Modules

7.1 LOAD_PATH (1.1)

Only modules located somewhere along the `LOAD_PATH` can be loaded with `using`.

But: If a directory contains `Project.toml`, it becomes a project directory and only entries listed in `Project.toml` can be loaded (even if the directory is on the `LOAD_PATH`).

As a general rule, though: If one has to fiddle with the `LOAD_PATH`, something is probably not right. Packages have all their dependencies in `Project.toml`. Anything that gets run from the `REPL` is `include`d. That only leaves potential startup code that sits in a `module` as a candidate for being on the `LOAD_PATH`.

7.2   Sub-Modules (1.1)

Functions from sub-modules can be exported by the main module. Example:

```
module scratch
export foo
module inner1
    export foo
    function foo()
        println("foo")
    end
end
using .inner1
end
```

7.3   Extending a function in another module (1.1)

The problem:

- Module `B` defines type `Tb` and function `foo(x ::  Tb)`.

- Module `A` contains a generic function `bar(x)` that calls `foo()`. It should use the `foo()` that matches the type of `x`. That is, when called as `foo(x ::  Tb)`, we want to call `B.foo`.

Solution:

- Module `A`:

– Define the stub: `function foo end`

– Call `foo(x)` from within `bar`.

- Module `B`:

  – Define `function foo(x :: Tb)`

  – `import A.foo`

- Now `A.bar(x)` knows about `B.foo()` and calls it when the type matches the signature.

See Duck typing when 'quack' is not in 'Base'.

8   Packages

8.1   Environments (1.1)

An environment is anything with a `Manifest.toml`. When you start Julia, you enter the version's environment (e.g. 1.1). When you add a package, you effectively edit `Manifest.toml`.

You can **add** additional environments using `Pkg.activate()` or `pkg> activate .` and then `Pkg.add` to initialize a `Mainfest.toml` in that directory. The environment determines how code is loaded.

- When you type `using M` Julia looks for module `M` in all directories that are listed in `LOAD_PATH`.

- Julia also looks in the directory of the currently activate package (which is **not** added to the `LOAD_PATH`). Exactly what `Pkg.activate()` does internally is not clear. Once you activate another package, previously activated packages are no longer considered during code loading.

- Note: Julia does **not** look in the current directory (unlike Matlab). In fact, the current directory really does nothing at all, except it is the base directory for `REPL` commands such as `cd()` or `include()`.

When examining a particular directory in `LOAD_PATH`, what happens depends on whether the directory contains `Manifest.toml` (or `Project.toml`; two go together).

- If it does not, Julia looks for `M.jl` in this directory.

- Otherwise, Julia **only** looks in `Manifest.toml`. The **only** part is key. Julia does not look in the directory itself.

8.1.1 Stacked environments

When you `activate` an environment, you do **not** deactivate previous environments. Instead, you now operate in a sort of union of all the environments that you activated during a session. This matters when both environments list the same packages in the Manifests.

Example: Start in environment `1.1` and `Pkg.add(D)`. `Pkg.activate(P)` and `Pkg.add(D)` with a different version of `D` (or using the local path for `D`). Which version of `D` is used after `using D`? The answer turns out to be that the most recent environment wins (that would usually be the currently activated project).

I encountered a case where I could not covince Julia to update an unregistered package, even using `Pkg.rm` followed by `Pkg.add`. The reason was that `1.1` referenced the same package, pointing to a fixed `github commit`.

8.2 Creating a package (1.1)

The easiest way is `PkgSkeleton.jl`. You need to set your `github` info (`user.name` etc) using

`git config --global user.name YourName`

This must be done insite a `git` directory. Then `generate` generates the directory structure and the required files (`Project.toml` etc). Example:

```
PkgSkeleton.generate("dir1/MyPackage")
```

Note: The package name should not end in ■`.jl`■ – this is automatically appended when the package is registered (?).

8.3 Package workflow (1.1)

Your packages will generally be unregistered. Your workflow needs to account for the fact that `Pkg` does not track versions for unregistered packages.

Here are the steps:

1. Initialize a `package` in a folder `pDir`; call the package P. This generates a directory structure with `src, test`, etc. If you plan on using this package as a dependency, it is best to place it in a sub-folder of `JULIA_PKG_DEVDIR` (`~/.julia/dev` by default). The reason is that `Pkg.develop` wants to download your code there.

2. While the code is being worked on: `Pkg.activate(ps)`. This makes sure that changes are written to the package's environment (`Project.toml`).

3. To add registered dependencies, simply use `Pkg.add(pkgName)`. No problem.

4. To add unregistered dependencies D that may change as you work on your project, use `Pkg.develop` instead.

   (a) Write code that makes a `PackageSpec` for D. This simplifies managing the package. Call this `ps`. `ps` should point to D's local directory, not to a `github` url. Otherwise, you end up tracking what is on github rather than your local edits.

   (b) `Pkg.develop(ps)` simply changes the entry for D in `Project.toml` from pointing at the `github repo` to pointing at the local dir. Key point: This is only operative while the environment P is active.

   (c) `Pkg.develop` is an alternative to `Pkg.add`, which edits `Project.toml` to point at `github`.

5. To freeze the state of the code:

   (a) `push P` and `D` to `github`.

   (b) in the environment for P: `Pkg.add(ps)` where `ps` should now point at the `github` url for D.

   (c) Even if you continue to push updates for unregistered dependencies to `github`, your package should track the fixed versions (identified by the `sha` key that defines the `commit`). Just don't run `Pkg.update`.

8.4   Unregistered packages as dependencies (1.1)

Important point: Unregistered packages need to be added as dependencies "by hand." `Pkg` cannot track when other packages depend on them. This is

a known issue 810. That means:

- Suppose you are working in `P` with dependency `D` that depends on `E`.

- `Pkg.add(D)` does not add `E` to P's `Project.toml`.

- You need to explicitly `Pkg.add(E)`.

Tracking changes in unregistered packages is difficult. While doing development work, adding dependencies with `Pkg.develop` seems to work. Once the code is copied to a remote computer, things get more complicated.

- The solution suggested on discourse suggests to always `develop` packages and to have relative paths in `Manifest.toml`. That would be relative paths of the form `../MyPackage`. User directory expansion, as in `~/abc` does not work.

Note:

- `Pkg.update` does nothing for unregistered dependencies.

- Deleting the corresponding subdirectory in `~/.julia/compiled` sometimes triggers a recompile, but not always.

- Removing a package and then adding it again sometimes triggers a recompile, but not always.

- `revise(MyPkg)` does not trigger a recompile.

It is easy enough to get the most recent version of the code on any computer. Forcing Julia to recompile once the code has changed is more difficult. On the local machine, `Revise` does the trick. But since you don't edit code on the remote machine before running it, I don't see how one can force a recompile.

What works sometimes:

- Restart the `REPL`.

8.5 Creating a package registry (1.1)

Any registry that lives in `~/.julia/registries` is automatically used by `Pkg`.

In principle, it is easy to create your own registry (see discourse for a guide). But there are problems:

1. Entries must be added to the registry by hand. Each package gets an entry line in the registry and a subdirectory in the registry directory with `Versions.toml, Deps.toml, Compat.toml, Package.toml`.

2. Any inconsistencies in the entries will cause the registry to be ignored by `Pkg`. So this approach is fragile.

3. Each time a package is changed or updated, the registry needs to be augmented by hand, including all dependencies. `Registrator.jl` is a project that aims to automate this, but the setup involved is substantial.

Fredrik Ekre has an example in his github repo. At this time, the approach is not really workable.

After trying a simple example registry with just one entry, I got an HTTP 404 error while adding a package. No info on which file was not accessible. This does not work.

8.6 Multiple Modules in one Package (1.1)

The cleanest approach is sub-modules. I.e.,

```
module Foo

include("That.jl")
include("Bar.jl")

using .Bar, .That

<code>

end
```

```
# In Bar.jl
module Bar
        using ..That
        <code>
end
```

One can still `import Foo.Bar` to only use the sub-module (especially for testing). In the test function, non-exported functions can be called as `Bar.f()`.

8.7   Testing a package (1.1)

Activate the package by issuing `activate .` in the package's directory (not in `src`). Then type `test`.

Note that the package needs the following in `Project.toml`:

```
[extras] Test = "8dfed614-e22c-5e08-85e1-65c5234f0b40"
[targets] test = ["Test"]
```

These are not automatically added. You need to hand-edit `Project.toml`. Or simply add `Test` as a dependency directly.

Placing test code inside a module:

- This can be useful when the test code defines `structs` that one would like to be able to modify without having to restart `Julia` all the time. Note that objects defined in tests are no longer visible once `Pkg` is exited.

- Place the module definition into `test`. Add `push(LOAD_PATH, @__DIR__)`. This has to be done in each module. Not elegant.

9   Operators

9.1   Logical (1.1)

`&&` is the logical AND operator, but in broadcasting use `.&` (even though `&` is a bitwise AND).

10   Performance

The compiler does not optimize out `if false` statements. Hence, defining a constant that switches self-testing code on and off does not result in no-ops. Of course, the overhead is quite small.

10.1   Profiling

The output generated by the built-in profiler is hard to read. `ProfileView` does not compile (1.1).

`StatProfilerHTML` is a good alternative (1.1). It provides a flame graph with clickable links that show which lines in a function take up most time.

10.2   Type stability

One can automate checking for type stability using the `code_warntype()` function. Example:

- For function `foo(x)`, call `code_warntype(stdout, foo, (Int,1))`.

- This can be written to a file by changing the `IO` argument.

- It generates output even if no issues are found.

- The amount of output generated is overwhelming. Signs of trouble are `Union` types, especially return types (at `Body:`).

11   Remote Clusters

11.1   Getting started with a test script

How to get your code to run on a typical Linux cluster?

- Get started by writing a simple test script (`Test3.jl`) so we can test running from the command line.

- Add the Julia binary to the `PATH` using (on Macos, editing `~/.bash_profile`):

  `PATH="/Applications/Julia-1.1.app/Contents/Resources/julia/bin:$PATH"`

- Then make sure you can run the test script with
  `julia ■/full/path/to/Test3.jl■`

Now copy `Test3.jl` to a directory on the cluster and repeat the same.

- You may need to add the Julia binary to the path. On Longleaf (editing `~/.bash_profile`):
  `export PATH="/nas/longleaf/apps/julia/1.1.0/bin:$PATH"`

- Then run `julia "/full/path/to/Test3.jl"`

Now run the test script via batch file:
```
sbatch -p general -N 1 -J "test_job" -t 3-00 --mem 16384 -n 1 --
mail-type=end --mail-user=lhendri@email.unc.edu -o "test1.out" -
-wrap="julia /full/path/to/Test3.jl"
```

11.2 Generate an ssh key

This allows log on without password. Instructions on the web.
Now you can use the terminal to log in with `ssh user@longleaf.unc.edu`.

11.3 Rsync File Transfer

A reliable command line transfer option is `rsync`. The command would be something like
```
rsync -atuzv "/someDirectory/sourceDir/" "username@longleaf.unc.edu:someDirectoryS
```
Notes:

- The source dir should end in "/"; the target dir should not.

- Exluding `.git` speeds up the transfer.

- `--delete` ensures that no old files remain on the server.

11.4 Git File Transfer

1. Change into the package directory (which is already a `git repo`).

2. Add a remote destination (once):
   `git remote add longleaf ssh://lhendri@longleaf.unc.edu/nas/longleaf/home/lhen`

3. Initialize the remote directory with a bare repo: `git init --bare`. Bare means that the actual files are not copied there. It needs to be bare so `push` does not produce errors later.

4. Verify the remote: `git remote show longleaf`

When files have changed:

1. Change into the package directory

2. `git commit -am ■commit message■`

3. `git push longleaf master`

Note that this does not upload any files! So this only works for packages, not for code that should be run outside of packages.

11.5 Running code on the cluster

Steps:

1. Copy your code and all of its dependencies to the cluster (see Section 11.3).

2. Write a Julia script that contains the startup code for the project and then runs the actual computation (call this `batch.jl`).

3. Write a batch file that submits `julia batch.jl` as a job to the cluster's job scheduler. For UNC's `longleaf` cluster, this would be slurm. So you need to write `job.sl` that will be submitted using `sbatch job.sl`.

11.5.1 The Julia script

Submitting a job is (almost) equivalent to `julia batch.jl` from the terminal.

- Note: `cd()` does not work in these command files. To include a file, provide a full path.

If you only use registered packages, life is easy. Your code would simply say:

```
using Pkg
# This needs to be run only once
Pkg.add(MyPackage)
# If you want the latest version each time
Pkg.update()
using MyPackage
MyPackage.run()
```

If `MyPackage` contains unregistered dependencies, things get more difficult. Now `batch.jl` must:

1. Activate the package's environment.

2. `develop` all unregistered dependencies. This replaces the invalid paths to directories on the local machine (e.g. `/Users/lutz/julia/...`) with the corresponding paths on the cluster (e.g. `/nas/longleaf/...`). Note: I verified that one cannot replace `homedir()` with ~ in Manifest.toml.

3. `using MyPackage`

4. `MyPackage.run()`

This approach requires you to keep track of all unregistered dependencies and where they are located on the remote machine. My way of doing this is contained in `PackageTools.jl` in the `shared` repo (this is not a package b/c its very purpose is to facilitate loading of unregistered packages).

For an example implementation of the entire process, see `batch_commands.jl` in `TestPkg2LH`.

- This uses `PackageToolsLH` to handle directories on different computers and file transfer.

- `write_command_file()` writes the julia file that is to be executed remotely (`command_file.jl`).

- `write_sbatch` writes the sbatch file that will be submitted to `slurm`.

- `project_upload()` uses `rsync` to copy the code of the project, its dependencies, and some general purpose code that is required at startup (mainly `PackageToolsLH` itself) to the remote machine.

11.5.2 The sbatch file

How this works can be looked up online. The only trick is that the Julia command requires a full path (or a relative path, but that's a little risky) on the remote machine.

`PackageToolsLH` keeps track of where things are on each machine. It is used to build the full paths.

11.5.3 Instantiating Packages

When packages are run, all dependencies must be installed. This would usually be done with `instantiate`. But this fails when the package is `develop`ed rather than `added`. Therefore: if a package fails to build or test (for example, after its first upload, or after new dependencies are installed that the remote machine does not have installed):

1. An indicator that a dependency is missing is the error message: `ERROR: MethodError:  Cannot 'convert' an object of type Nothing to an object of type Base.SHA1`

2. Switch to a test environment where one can mess up the `Project.toml`.

3. `Pkg.add(ps)` where `ps` is the `PackageSpec` for the package that does not build. It must point at the `github` url.

4. This is not always enough. In that case, `activate` the package that does not build. Use `>pkg st -m` to show the packages that are not loaded and simply `add` them until the package builds and tests.

Now the package can be built or developed everywhere.

Sometimes old versions of `Project.toml` lie around somewhere (where?) in the Julia installation. They may contain dependencies that don't exist anymore. Then the package does not build. The only solution that seems to work: `Pkg.add` the package from somewhere with a `PackageSpec` that points at `github`.

- For this purpose, it is useful to have an environment lying around that is just for adding packages that need to be downloaded.

12 Types (1.1)

I find it easiest to write model specific code NOT using parametric types. Instead, I define `type aliases` for the types used in custom types (e.g., `Double=Float64`). Then I hardwire the use of `Double` everywhere. This removes two problems:

1. Possible type instability as the compiler tries to figure out the types of the custom type fields.

2. It becomes possible to call constructors with, say, integers of all kinds without raising method errors.

12.1 Constructors (1.1)

Constructing objects with many fields:

- Define an inner constructor that leaves the object (partially) uninitialized. It is legal to have new(x) even if the object contains additional fields.

12.2 Inheritance (1.1)

There is no inheritance in Julia. Abstract types have no fields and concrete types have no subtypes.

There are various discussions about how to implement types that share common fields.

For simple cases, it is probably best to just repeat the fields in all types. This can be automated using `@forward` in `Lazy.jl`.

One good piece of advice: ensure that methods are generally defined on the abstract type, so that all concrete types have the same interface (kind of the point of having an abstract type).

12.3 Loading and saving (1.1)

`using FileIO` and extension `.jld2` automatically saves in `jld2` format. This can save used defined types.

Loading user defined types is more complicated. All modules needed to construct the loaded types need to be known in the loading module and in `Main`. See Issue 134. It is not possible to use `Core.eval(Main, :(using Module))` for unclear reasons.

Implications:

1. Each user defined type needs its own `load` function.

2. All dependencies need to imported into Main ∎by hand∎ `for each loaded object`.

An alternative is `BSON.jl`. It has the same limitation.

One could save the `ParamVector`s in each object and reconstruct the object from those (recursively). This, of course, only works for objects that can be constructed from `ParamVector`s. Each `ParamVector` could be stored as a `Dict{Symbol, Any}`. But even easier: store the `ParamVector`s directly. Constructing them after loading only requires `modelLH`. The approach would then be:

1. Collect the `ParamVector`s from all model objects into a `Dict{Symbol,ParamVector}`. The symbol identifies the associated model object.

2. Save the `Dict`.

3. In `Main:` using modelLH, so that loading works.

4. Function that loads the model:

    (a) Construct the model object with arbitrary default values.

(b) Load the `ParamVector`s.

(c) Sync each `ParamVector`'s parameters into the correct model object. Essentially, the model object needs a constructor that accepts a `ParamVector`.

13   Unit Testing (1.1)

All code should be in `modules` because code in `Main` runs slower, pollutes `Main`, and it harder to `revise`. This also applies to test code.

However, placing the `@test` or `@testset` portions into the test module causes them not to run sometimes (why?). It also implies that `using` the test module runs all tests, which is generally unwanted. I therefore place the `@test` code into a separate file (not inside a `module`).

Errors in the code to be tested (but not caught by `@test`) cause the entire test run to crash. Preventing this requires all tests to be enclosed in a `@testset`. A sequence of `@testset` does not do the trick. An error in one prevents all others from being run. Nested `@testset`s produce nested error reports (nice).

`@test` statements can be placed inside functions. To preserve result reporting, the function should contain a @testset and return its result.

To run tests selectively, simply `include` the file that contains the `@testset` at the `REPL`. This, of course, requires that tests are organized into separate files.

14   Workflow (1.1)

`Revise` is key. It is now possible to simply use `using` on any `module` once. `Revise` then automatically keeps track of changes. Using `includet` creates problems for me.

But keep in mind that `Revise` cannot handle:

1. changes in file structure (you factor out some code into a new file that is `include`d in the main file);

2. changes in `struct`s.

Those still require restarting the REPL.